

# Zhiheng Li

🏠: <https://zhiheng.li>

✉: [zhiheng.li@ieee.org](mailto:zhiheng.li@ieee.org)

📞: +1-585-351-1179

🌐: <https://www.linkedin.com/in/zhiheng-li-ur/>

🐙: <https://github.com/zhihengli-UR>

## EDUCATION

### University of Rochester

Ph.D. in Computer Science

Advisor: Chenliang Xu

Rochester, NY, USA

Aug 2018 - Jul 2023

### Wuhan University

B.Eng. in Software Engineering

Wuhan, Hubei, China

Sep 2014 - June 2018

## RESEARCH INTERESTS

Computer Vision, Data-centric Machine Learning, Generative AI, Responsible AI

## WORKING EXPERIENCE

### Amazon AGI (Olympus Multimodal Pretraining team)

Bellevue, WA, USA

Applied Scientist II

Nov 2023 - Present

Project: **Pretraining Data for Amazon Nova Reel**—Amazon’s Video Generation Foundation Model

- Built scalable PySpark-based data pipeline to filter, enrich, and normalize large-scale video generation pretraining data.
- Created comprehensive data annotation guidelines and supervised data labeling process for dense synthetic video captioning; validated improved video-text alignment performance through data ablation.
- Developed a data-centric method for camera movement control via prompting (see more in camera control prompting guide); Achieved near-perfect results on VBench’s camera motion control metric.
- Collaborated with Responsible AI (RAI), engineering, and legal teams to filter pre-training data for the video generation model, ensuring full compliance with RAI and legal requirements.
- Helped with building the data pipeline for talking head video generation.

### AWS Bedrock (Bedrock Multimodal Responsible AI team)

Bellevue, WA, USA

Applied Scientist II

Jul 2023 - Oct 2023

Project: **Fairness for vision-language foundation models**

- Developed a model to address the fairness-performance trade-off.
- Created comprehensive data annotation guidelines and supervised data labeling process.
- Conducted performance and fairness evaluation of the designed solution for multimodal models.
- Implemented the fairness solution into production for Titan Multimodal models.

### Meta AI (Responsible AI: Robustness and Safety Team)

Seattle, WA, USA

Research Intern

May 2022 - Dec 2022

Project: **Multi-shortcut Mitigation**

Mentors: Ivan Evtimov, Mark Ibrahim, Albert Gordo

- **Problem Setup:** While previous works have the tenuous assumption that only a single shortcut (i.e., bias) exists, images are rife with multiple visual cues. Thus, we study a neglected but critical problem—multi-shortcut learning in image classification.
- **New Datasets:** We create two datasets—(1) *UrbanCars*, a semi-synthetic dataset with multiple precisely controlled spurious cues; (2) *ImageNet-W*, a new out-of-distribution variant of ImageNet to evaluate the reliance on *watermark*—a new shortcut that we discovered in ImageNet, which affects a broad range of vision models, including supervised ResNet, self-supervised models (e.g., MoCoV3 and MAE), and even large foundation models (e.g., SEER, SWAG, and CLIP).
- **New Finding:** Benchmarking on both datasets reveals the overlooked challenge in shortcut learning that resembles a Whac-A-Mole game—mitigating one shortcut amplifies others.
- **Method:** We propose Last Layer Ensemble (LLE) based on data augmentation to mitigate multiple shortcuts jointly.
- **Results:** Our proposed LLE method does not suffer the Whac-A-Mole problem.
- **Publication:** The paper [1] is accepted to CVPR 2023.

### NEC Labs America

Princeton, NJ, USA

Research Intern

Jun 2021 - Sep 2021

Project: **Compositional Text-to-Image Synthesis**

Mentors: Martin Renqiang Min, Kai Li

- **Problem Setup:** Existing text-to-image methods lack compositionality. E.g., They struggle to generate images of “*a man wearing lipstick*,” leading to fairness and robustness issues.

- **Method:** We propose StyleT2I, a novel text-to-image synthesis framework. To improve compositionality, we propose: (1) CLIP-guided contrastive loss; (2) leveraging disentangled representations in StyleGAN to control image generation.
- **Results:** StyleT2I outperforms previous methods, e.g., ControlGAN and TediGAN, on CelebA-HQ and CUB datasets.
- **Publication:** The paper [2] is accepted to CVPR 2022.

---

## RESEARCH PROJECTS AT UNIVERSITY OF ROCHESTER

---

### Project: **Unknown Bias Discovery and Mitigation**

Apr 2021 - Feb 2022

- **Problem Setup:** Biases in datasets are usually unknown in terms of (1) bias type; and (2) number of biases. Even for known bias, bias labels (e.g., gender labels) are hard to collect due to ethical and privacy concerns. The goal is to discover and mitigate unknown biases without any prior knowledge of bias types or numbers of biases.
- **Method:** We propose DebiAN, which jointly trains two networks, i.e., a classifier and a discoverer, in an alternate fashion to discover and mitigate unknown biases.
- **New Dataset:** We construct Multi-Color MNIST, a new dataset containing multiple biases, to study the mitigation of multiple biases, which is underexplored in works.
- **Results:** DebiAN outperforms existing approaches, e.g., LfF and EIIL, on CelebA, BAR, bFFHQ, and Places datasets. Besides, our Multi-Color MNIST dataset demonstrates DebiAN's advantage of mitigating multiple biases, whereas LfF, a previous seminal work, can only identify and mitigate one bias. Besides, DebiAN discovers interesting unknown biases, e.g., visible hair area for gender classification and indoor vs. outdoor scene for restaurant class in scene classification.
- **Publication:** The paper [3] is accepted to ECCV 2022.

### Project: **Unknown Bias Discovery**

Sep 2020 - Mar 2021

- **Problem Setup:** Existing bias detection frameworks heavily rely on human assumptions of the bias type. As a result, unknown bias cannot be discovered, not to mention to be mitigated. Thus, we formalize a new task, Unknown Bias Discovery.
- **Method:** Leveraging the disentangled latent space of generative models (e.g., StyleGAN), we optimize the latent hyperplane that represents an unknown bias.
- **Results:** Our method discovers some interesting unknown biases in ImageNet dataset (e.g., the shade of fur color for cat) and Places dataset (e.g., Eiffel tower for towel class in scene classification).
- **Publication:** The paper [4] is accepted to ICCV 2021.

### Project: **Interpretable Visual Representation via Perceptual Grouping on CNN**

Jun 2019 - Mar 2020

- **Problem Setup:** The commonly used convolutional neural nets (CNN) fall short in interpretability and efficiency. We address the limitations by incorporating the traditional perceptual grouping process into CNNs.
- **Method:** We propose Deep Group Model (DGM), which performs grouping hierarchically over CNN feature maps.
- **Results:** DGM outperforms existing approaches in unified perceptual parsing task, which consists of image-level scene classification and multiple pixel-level segmentation tasks, e.g., object segmentation, part segmentation, etc. Besides, DGM reduces computational overhead compared to pixel-based self-attention approaches. Moreover, DGM facilitates explainable CNN by generating saliency maps with sharp boundaries that locate spatial regions more precisely.
- **Publication:** The paper [5] is accepted to CVPR 2020.

### Project: **Lip Movements Generation**

Jun 2017 - Mar 2018

- **Problem Setup:** Given audio and an input lip image, the task is to generate a lip movements video of the mouth speaking the voice in the given audio.
- **Method:** Based on conditional GAN, we propose novel losses to model the correlation between lip movements and audio.
- **Results:** The generated videos have good visual quality. The generated lip movements and audio are in good synchronization, facilitating many future works on the talking head generation task.
- **Publication:** The paper [6] is accepted to ECCV 2018.

---

## PUBLICATION

---

- [1] **Zhiheng Li**, Ivan Evtimov, Albert Gordo, Caner Hazirbas, Tal Hassner, Cristian Canton Ferrer, Chenliang Xu, and Mark Ibrahim. A Whac-A-Mole Dilemma: Shortcuts Come in Multiples Where Mitigating One Amplifies Others. In *CVPR*, 2023.
- [2] **Zhiheng Li**, Martin Renqiang Min, Kai Li, and Chenliang Xu. StyleT2I: Toward Compositional and High-Fidelity Text-to-Image Synthesis. In *CVPR*, 2022.
- [3] **Zhiheng Li**, Anthony Hoogs, and Chenliang Xu. Discover and Mitigate Unknown Biases with Debiasing Alternate Networks. In *ECCV*, 2022.

- [4] **Zhiheng Li** and Chenliang Xu. Discover the Unknown Biased Attribute of an Image Classifier. In *ICCV*, 2021.
- [5] **Zhiheng Li**, Wenxuan Bao, Jiayang Zheng, and Chenliang Xu. Deep Grouping Model for Unified Perceptual Parsing. In *CVPR*, 2020.
- [6] Lele Chen\*, **Zhiheng Li\*** (**\*Equal Contribution**), Ross K Maddox, Zhiyao Duan, and Chenliang Xu. Lip Movements Generation at a Glance. In *ECCV*, 2018.
- [7] Raghav Mehta, Vitor Albiero, Li Chen, Ivan Evtimov, Tamar Glaser, **Zhiheng Li**, and Tal Hassner. You Only Need a Good Embeddings Extractor to Fix Spurious Correlations. In *ECCV Responsible Computer Vision Workshop*, 2022.
- [8] Jie Chen, **Zhiheng Li**, Jiebo Luo, and Chenliang Xu. Learning a Weakly-Supervised Video Actor-Action Segmentation Model with a Wise Selection. In *CVPR*, 2020.
- [9] **Zhiheng Li**, Geemi P. Wellawatte, Maghesree Chakraborty, Heta A. Gandhi, Chenliang Xu, and Andrew D. White. Graph Neural Network Based Coarse-Grained Mapping Prediction. *Chemical Science*, 2020.

## PROFESSIONAL SERVICES

---

- **Conference Reviewer:** NeurIPS'20, CVPR'21, ICML'21, ICCV'21, NeurIPS'21, ICLR'22, AAAI'22, CVPR'22, ICML'22, ECCV'22, NeurIPS'22 (top reviewer), CVPR'23, FAccT'23, ICML'23, ICCV'23, NeurIPS'23, ICLR'24, CVPR'24, ICML'24, ECCV'24
- **Jornal Reviewer:** TMLR, TPAMI, TMM
- **Volunteer:** FAccT'21, ICLR'21

## SKILLS

---

Python, PyTorch, PySpark, Git